

Elias Stengel-Eskin

Website: esteng.github.io
Email: esteng@cs.utexas.edu
LinkedIn: [elias-stengel-eskin](https://www.linkedin.com/in/elias-stengel-eskin/)
GitHub: github.com/esteng

EDUCATION

Johns Hopkins University

Ph.D. in Computer Science – Advisor: *Benjamin Van Durme*

Baltimore, USA

2018–2023

- Ph.D. thesis: “Modeling Meaning for Description and Interaction”
- supported by NSF Graduate Research Fellowship

Johns Hopkins University

MSE in Computer Science – Advisor: *Benjamin Van Durme*

Baltimore, USA

2018–2021

McGill University

Bachelor of Arts and Sciences in Cognitive Science

Montréal, Canada

2014–2018

- Minor: Linguistics
- First Class Honours
- Honours thesis: “Variational Bayesian Inference for Unsupervised Lexicon Discovery” – Advisor: *Timothy O’Donnell*

EXPERIENCE

The University of Texas at Austin

Assistant Professor, Department of Computer Science

August 2025–Present

Microsoft

Researcher (Consulting)

February 2025–Present

University of North Carolina, Chapel Hill

Postdoctoral Research Associate – Advisor: *Mohit Bansal*

August 2023–August 2025

Microsoft Research

PhD Research Intern – Advisors: *Marc-Alexandre Côté, Eric Yuan, Pierre-Yves Oudeyer*

March 2022–March 2023

Microsoft Research - Semantic Machines

PhD Research Intern – Advisor: *Yu Su*

May 2021 – August 2021

Montreal Computational and Quantitative Linguistics Lab

Research Assistant – Advisor: *Morgan Sonderegger*

2016–2018

PUBLICATIONS (PEER-REVIEWED)

- [1] Jio Choi, Mohit Bansal, and **Elias Stengel-Eskin**, “Language models identify ambiguities and exploit loopholes”, *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2025.
- [2] Justin Chih-Yao Chen, Archiki Prasad, Swarnadeep Saha, **Elias Stengel-Eskin**, and Mohit Bansal, “Magicore: Multi-agent, iterative, coarse-to-fine refinement for reasoning”, *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2025.

- [3] Duy Nguyen, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “Laser: Learning to adaptively select reward models with multi-armed bandits”, *Advances in Neural Information Processing Systems 39*, 2025.
- [4] Eduardo Calò, Guanyi Chen, **Elias Stengel-Eskin**, Albert Gatt, and Kees van Deemter, “Incorporating formulaicness in the automatic evaluation of naturalness: A case study in logic-to-text generation”, in *Proceedings of the 18th International Natural Language Generation Conference*, 2025, pages 352–365.
- [5] Nithin Sivakumaran, Chia-Yu Yang, Abhay Zala, Shoubin Yu, Daeun Hong, Xiaotian Zou, **Elias Stengel-Eskin**, Dan Carpenter, Wookhee Min, Cindy Hmelo-Silver, *et al.*, “A multimodal classroom video question-answering framework for automated understanding of collaborative learning”, in *Proceedings of the 27th International Conference on Multimodal Interaction*, 2025, pages 516–525.
- [6] Han Wang, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “Retrieval-augmented generation with conflicting evidence”, *Conference on Language Modeling (COLM)*, 2025.
- [7] Hanqi Xiao, Yi-Lin Sung, **Elias Stengel-Eskin**, and Mohit Bansal, “Task-circuit quantization: Leveraging knowledge localization and interpretability for compression”, *Conference on Language Modeling (COLM)*, 2025.
- [8] Archiki Prasad, **Elias Stengel-Eskin**, Justin Chih-Yao Chen, Zaid Khan, and Mohit Bansal, “Learning to generate unit tests for automated debugging”, *Conference on Language Modeling (COLM)*, 2025.
- [9] David Wan, Eran Hirsch, **Elias Stengel-Eskin**, Ido Dagan, and Mohit Bansal, “Generationprograms: Fine-grained attribution with executable programs”, *Conference on Language Modeling (COLM)*, 2025.
- [10] Atin Pothiraj, **Elias Stengel-Eskin**, Jaemin Cho, and Mohit Bansal, “Capture: Evaluating spatial reasoning in vision language models via occluded object counting”, *Proceedings of the International Conference on Computer Vision, (ICCV)*, 2025.
- [11] Duy Nguyen, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “Multi-attribute steering of language models via targeted intervention”, *Proceedings of the Sixty-Third Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025.
- [12] Eran Hirsch, Aviv Slobodkin, David Wan, **Elias Stengel-Eskin**, Mohit Bansal, and Ido Dagan, “Laquer: Localized attribution queries in content-grounded generation”, *Proceedings of the Sixty-Third Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025.
- [13] Ziyang Wang*, Shoubin Yu*, **Elias Stengel-Eskin***, Jaehong Yoon, Feng Cheng, Gedas Bertasius, and Mohit Bansal, “VideoTree: Adaptive tree-based video representation for LLM reasoning on long videos”, *Proceedings of the Forty-Second Annual IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
- [14] Zaid Khan, **Elias Stengel-Eskin**, Jaemin Cho, and Mohit Bansal, “DataEnvGym: Data generation agents in teacher environments with student feedback”, *The Thirteenth International Conference on Learning Representations (ICLR) Spotlight*, 2025.
- [15] Swarnadeep Saha, Archiki Prasad, Justin Chih-Yao Chen, Peter Hase, **Elias Stengel-Eskin**, and Mohit Bansal, “System-1.x: Learning to balance fast and slow planning with language models”, *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- [16] Amith Ananthram, **Elias Stengel-Eskin**, Carl Vondrick, Mohit Bansal, and Kathleen McKeown, “See it from my perspective: Diagnosing the western cultural bias of large vision-language models in image understanding”, *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- [17] **Elias Stengel-Eskin**, Peter Hase, and Mohit Bansal, “Teaching models to balance resisting and accepting persuasion”, *Proceedings of the Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)*, 2025.

- [18] David Wan, Justin Chih-Yao Chen, **Elias Stengel-Eskin**, and Mohit Bansal, “MAMM-Refine: A recipe for improving faithfulness in generation with multi-agent collaboration”, *Proceedings of the Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)*, 2025.
- [19] Han Wang, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “AdaCAD: Adaptively decoding to balance conflicts between contextual and parametric knowledge”, *Proceedings of the Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)*, 2025.
- [20] **Elias Stengel-Eskin**, Peter Hase, and Mohit Bansal, “LACIE: Listener-aware finetuning for confidence calibration in large language models”, *Advances in Neural Information Processing Systems 38*, 2024.
- [21] Jinhao Duan, Renming Zhang, James Diffenderfer, Bhavya Kailkhura, Lichao Sun, **Elias Stengel-Eskin**, Mohit Bansal, Tianlong Chen, and Kaidi Xu, “GTBench: Uncovering the strategic reasoning limitations of LLMs via game-theoretic evaluations”, *Advances in Neural Information Processing Systems 38*, 2024.
- [22] Ansel Blume, Khanh Duy Nguyen, Zhenhailong Wang, Yangyi Chen, Michal Shlapentokh-Rothman, Xiaomeng Jin, Jeonghwan Kim, ..., **Stengel-Eskin-Elias**, et al., “MIRACLE: An online, explainable multimodal interactive concept learning system”, in *Proceedings of the 32nd ACM International Conference on Multimedia*, 2024.
- [23] Peter Hase, Thomas Hofweber, Xiang Zhou, **Elias Stengel-Eskin**, and Mohit Bansal, “Fundamental problems with model editing: How should rational belief revision work in LLMs?”, *Transactions on Machine Learning*, 2024.
- [24] **Elias Stengel-Eskin**, Archiki Prasad, and Mohit Bansal, “ReGAL: Refactoring programs to discover generalizable abstractions”, *The 41st International Conference on Machine Learning (ICML)*, 2024.
- [25] Justin Chih-Yao Chen, Swarnadeep Saha, **Elias Stengel-Eskin**, and Mohit Bansal, “MAGDi: Structured distillation of multi-agent interaction graphs improves reasoning in smaller language models”, *The 41st International Conference on Machine Learning (ICML)*, 2024.
- [26] Haotian Fu, Pratyusha Sharma, **Elias Stengel-Eskin**, George Konidaris, Nicolas Le Roux, Marc-Alexandre Côté, and Xingdi Yuan, “Language-guided skill learning with temporal variational inference”, *The 41st International Conference on Machine Learning (ICML)*, 2024.
- [27] David Wan, Jaemin Cho, **Elias Stengel-Eskin**, and Mohit Bansal, “Contrastive region guidance: Improving grounding in vision-language models without training”, *The 18th European Conference on Computer Vision (ECCV)*, 2024.
- [28] Han Wang*, Archiki Prasad*, **Elias Stengel-Eskin***, and Mohit Bansal, “Soft self-consistency improves language model agents”, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- [29] Maryam Hashemzadeh, **Elias Stengel-Eskin**, Sarath Chandar, and Marc-Alexandre Cote, “Sub-goal distillation: A method to improve small language agents”, *The 3rd Conference on Lifelong Learning Agents (CoLLAs)*, 2024.
- [30] **Elias Stengel-Eskin**, Kyle Rawlins, and Benjamin Van Durme, “Zero and few-shot semantic parsing with ambiguous inputs”, in *The Twelfth International Conference on Learning Representations (ICLR)*, 2024.
- [31] Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “Rephrase, augment, reason: Visual grounding of questions for vision-language models”, in *The Twelfth International Conference on Learning Representations (ICLR)*, 2024.
- [32] **Elias Stengel-Eskin** and Benjamin Van Durme, “Did You Mean...? Confidence-based Trade-offs in Semantic Parsing”, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.

- [33] Elias Stengel-Eskin and Benjamin Van Durme, “Calibrated Interpretation: Confidence Estimation in Semantic parsing”, *Transactions of the Association for Computational Linguistics (TACL)*, 2023.
- [34] Elias Stengel-Eskin, Jimena Guallar-Blasco, Yi Zhou, and Benjamin Van Durme, “Why Did the Chicken Cross the Road? Rephrasing and Analyzing Ambiguous Questions in VQA”, *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023.
- [35] Zhuowan Li, Xingrui Wang, Elias Stengel-Eskin, Adam Kortylewski, Wufei Ma, Benjamin Van Durme, and Alan Yuille, “Super-CLEVR: A Virtual Benchmark to Diagnose Domain Robustness in Visual Reasoning”, *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [36] Elias Stengel-Eskin and Benjamin Van Durme, “The Curious Case of Control”, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022.
- [37] Elias Stengel-Eskin, Emmanouil Antonios Platanios, Adam Pauls, Sam Thomson, Hao Fang, Benjamin Van Durme, Jason Eisner, and Yu Su, “When More Data Hurts: A Troubling Quirk in Developing Broad-Coverage Natural Language Understanding Systems”, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022.
- [38] Chenyu Zhang, Benjamin Van Durme, Zhuowan Li, and Elias Stengel-Eskin, “Visual Commonsense in Pretrained Unimodal and Multimodal Models”, in *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, Seattle, Washington: Association for Computational Linguistics, 2022.
- [39] Elias Stengel-Eskin, Andrew Hundt, Zuhong He, Aditya Murali, Nakul Gopalan, Matthew Gombolay, and Gregory D. Hager, “Guiding Multi-Step Rearrangement Tasks with Natural Language Instructions”, in *5th Annual Conference on Robot Learning (CoRL)*, 2021.
- [40] Zhuowan Li, Elias Stengel-Eskin, Yixiao Zhang, Cihang Xie, Quan Tran, Benjamin Van Durme, and Alan Yuille, “Calibrating Concepts and Operations: Towards Symbolic Reasoning on Real Images”, in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021.
- [41] Elias Stengel-Eskin, Jimena Guallar-Blasco, and Benjamin Van Durme, “Human-Model Divergence in the Handling of Vagueness”, in *Proceedings of the 1st Workshop on Understanding Implicit and Underspecified Language (UnImplicit)*, Online: Association for Computational Linguistics, 2021.
- [42] Elias Stengel-Eskin, Kenton Murray, Sheng Zhang, Aaron Steven White, and Benjamin Van Durme, “Joint Universal Syntactic and Semantic Parsing”, *Transactions of the Association for Computational Linguistics (TACL)*, 2021.
- [43] Elias Stengel-Eskin, Jimena Guallar-Blasco, and Benjamin Van Durme, “Exploring Human-Model Divergence Through Vagueness”, *Proceedings of the Society for Computation in Linguistics (SCiL)*, 2021, *Abstract.
- [44] Ryan Culkin, J. Edward Hu, Elias Stengel-Eskin, Guanghui Qin, and Benjamin Van Durme, “Iterative Paraphrastic Augmentation with Discriminative Span Alignment”, *Transactions of the Association for Computational Linguistics (TACL)*, volume 9, 2021.
- [45] Elias Stengel-Eskin, Aaron Steven White, Sheng Zhang, and Benjamin Van Durme, “Universal Decompositional Semantic Parsing”, in *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2020.
- [46] Aaron Steven White, Elias Stengel-Eskin, Siddharth Vashishtha, Venkata Subrahmanyam Govindarajan, Dee Ann Reisinger, Tim Vieira, Keisuke Sakaguchi, Sheng Zhang, Francis Ferraro, Rachel Rudinger, et al., “The Universal Decompositional Semantics Dataset and Decomp Toolkit”, in *Proceedings of The 12th Language Resources and Evaluation Conference*, 2020.

- [47] Elias Stengel-Eskin, Tzu-Ray Su, Matt Post, and Benjamin Van Durme, “A Discriminative Neural Model for Cross-Lingual Word Alignment”, in *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, 2019.
- [48] Michael McAuliffe, Elias Stengel-Eskin, Michaela Socolof, and Morgan Sonderegger, “Polyglot and Speech Corpus Tools: A System for Representing, Integrating, and Querying Speech Corpora.” In *INTERSPEECH*, 2017.

PUBLICATIONS (OTHER)

1. Amith Ananthram, Elias Stengel-Eskin, Lorena A Bradford, Julia Demarest, Adam Purvis, Keith Krut, Robert Stein, Rina Elster Pantalony, Mohit Bansal, Kathleen McKeown, “PoSh: Using Scene Graphs To Guide LLMs-as-a-Judge For Detailed Image Descriptions” (2025), *Under Review*
2. Zaid Khan, Archiki Prasad, Elias Stengel-Eskin, Jaemin Cho, Mohit Bansal, “One Life to Learn: Inferring Symbolic World Models for Stochastic Environments from Unguided Exploration” (2025), *Under Review*
3. Jaewoo Lee, Archiki Prasad, Justin Chih-Yao Chen, Zaid Khan, Elias Stengel-Eskin, Mohit Bansal, “PRInTS: Reward Modeling for Long-Horizon Information Seeking” (2025), *Under Review*
4. Joykirat Singh, Justin Chih-Yao Chen, Archiki Prasad, Elias Stengel-Eskin, Akshay Nambi, Mohit Bansal, “Think Right: Learning to Mitigate Under-Over Thinking via Adaptive, Attentive Compression” (2025), *Under Review*
5. Victor Wang, Elias Stengel-Eskin, “Calibrating Verbalized Confidence with Self-Generated Distractors” (2025), *Under Review*
6. Vaidehi Patil, Elias Stengel-Eskin, Mohit Bansal, “The Sum Leaks More Than Its Parts: Compositional Privacy Risks and Mitigations in Multi-Agent Collaboration” (2025), *Under Review*
7. Hanqi Xiao, Vaidehi Patil, Hyunji Lee, Elias Stengel-Eskin, Mohit Bansal, “Generalized Correctness Models: Learning Calibrated and Model-Agnostic Correctness Predictors from Historical Patterns” (2025), *Under Review*
8. Tianyi Niu, Jaemin Cho, Elias Stengel-Eskin, Mohit Bansal, “RotBench: Evaluating Multimodal Large Language Models on Identifying Image Rotation” (2025), *Under Review*
9. Duy Nguyen, Archiki Prasad, Elias Stengel-Eskin, Mohit Bansal, “GrAInS: Gradient-based Attribution for Inference-Time Steering of LLMs and VLMs” (2025), *Under Review*
10. Hyunji Lee, Seunghyun Yoon, Yunjae Won, Hanseok Oh, Geewook Kim, Trung Bui, Franck Dernoncourt, Elias Stengel-Eskin, Mohit Bansal, Minjoon Seo, “Context-Informed Grounding Supervision” (2025), *Under Review*
11. David Wan, Han Wang, Elias Stengel-Eskin, Jaemin Cho, Mohit Bansal, “CLaMR: Contextualized Late-Interaction for Multimodal Content Retrieval” (2025), *Under Review*
12. Emmanouil Zaranis, António Farinhas, Saul Santos, Beatriz Canaverde, Miguel Moura Ramos, Aditya K Surikuchi, André Viveiros, Baohao Liao, Elena Bueno-Benito,..., Nithin Sivakumaran, Elias Stengel-Eskin, et al., “Movie Facts and Fibs (MF²): A Benchmark for Long Movie Understanding” (2025), *Under Review*
13. Ron Eliav, Arie Cattan, Eran Hirsch, Shahaf Bassan, Elias Stengel-Eskin, Mohit Bansal, Ido Dagan, “CLATTER: Comprehensive Entailment Reasoning for Hallucination Detection” (2025), *Under Review*
14. Justin Chih-Yao Chen, Sukwon Yun, Elias Stengel-Eskin, Tianlong Chen, Mohit Bansal, “Symbolic mixture-of-experts: Adaptive skill-based routing for heterogeneous reasoning” (2025), *Under Review*
15. Zaid Khan, Elias Stengel-Eskin, Archiki Prasad, Jaemin Cho, Mohit Bansal, “Executable Functional Abstractions: Inferring Generative Programs for Advanced Math Problems” (2025), *Under Review*
16. Vaidehi Patil, Elias Stengel-Eskin, Mohit Bansal, “UPCORE: Utility-Preserving Coreset Selection for Balanced Unlearning” (2025), *Under Review*

17. Yue Huang, Chujie Gao, Siyuan Wu, ..., **Elias Stengel-Eskin**, ..., Bo Li, Xiangliang Zhang, “On the Trustworthiness of Generative Foundation Models: Guideline, Assessment, and Perspective” (2025), *Preprint*
18. Thomas Hofweber, Peter Hase, **Elias Stengel-Eskin**, Mohit Bansal, “Are language models rational? The case of coherence norms and belief revision” (2024), *Preprint*
19. **Elias Stengel-Eskin**, “Modeling Meaning for Description and Interaction”, (2023) **Ph.D. Thesis**
20. Jimena Guallar-Blasco, **Elias Stengel-Eskin**, Benjamin Van Durme, “Analyzing Question Ambiguity in *Why*-Questions” (2023), *Mid-Atlantic Student Colloquium on Speech, Language and Learning (MASC-SLL 2023)*
*Abstract
21. Shalaka Vaidya, **Elias Stengel-Eskin**, Joao Sedoc, “Automatic Evaluation of Chit-chat via Semantic Parsing” (2022), *Mid-Atlantic Student Colloquium on Speech, Language and Learning (MASC-SLL 2022)**Abstract
22. Yunmo Chen, Seth Ebner, Tongfei Chen, Patrick Xia, **Elias Stengel-Eskin**, Tzu-Ray Su, J. Edward Hu, Nils Holzenberger, Ryan Culkin, Craig Harman, Max Thomas, Thomas Lippincott, Aaron Steven White, Kyle Rawlins, Benjamin Van Durme, “NIST TAC SM-KBP 2019 System Description: JHU/UR Framework”, (2019)
23. **Elias Stengel-Eskin**, “Variational Bayesian Inference for Unsupervised Lexicon Discovery” (2017), **Undergraduate Honours Thesis**

FUNDING

Received

- **IARPA BENGAL** 2025-2027
\$290k. Role: Postdoc
- **NSF NAIRR Pilot Grant** 2024-2025
Supervising Calibration for Safe and Trustworthy Artificial Intelligence Systems
2,500 GPU Hours. Role: **PI**
- **Cisco Gift** 2025
Improving LLM Reliability and Trustworthiness via Overconfidence Reduction and Conversation
\$75K. Role: Co-PI
- **CapitalOne Gift** 2025
Multi-Agent Systems and Mixtures-of-Experts for Robust Generation
Role: Co-PI
- **Center for AI Safety Compute Grant** 2024
Listener-Aware Confidence Calibration for Improving LLM Reliability and Trustworthiness
Unlimited A100 GPU access. Role: Co-PI
- **NSF Graduate Research Fellowship** 2018-2023
Role: Student

NEWS COVERAGE

1. *The importance of ambiguity*, Johns Hopkins University CS News, 04/15/2024
2. *Speaking easy: A simple method to clarify voice commands*, Johns Hopkins University CS News, 01/08/2024

INVITED TALKS/PANELS

- **Towards Granular, Robust, Calibrated Judgments via Skill-Based and Persuasion-Balanced Reasoning** December 2, 2025
NeurIPS Expo Workshop on Exploring Trust and Reliability in LLM Evaluation
- **Panel:** Agentic AI in Wireless Networking and Communications November 14, 2025
UT Austin Wireless Networking and Communications Group
- **Panel:** The true impact of AI on the labor market September 30, 2025
Hire EZ
- **Communication, Calibration, and Grounding for Collaborative AI Agents** January 8, 2025
Duke University
- **Confidence-based Rephrasing, Refinement, and Selection** March 22nd, 2024
Keynote – UncertaiNLP Workshop, EACL 2024
- **Ambiguity in NLP** October 26th, 2023
UNC Chapel Hill – Linguistics Department
- **Uncertainty and Ambiguity in Semantic Structures** October 6th, 2023
Emory University
- **Language, Structure, and Beyond** May 9th, 2023
UNC Chapel Hill – MURGe Group
- **Language, Structure, and Beyond** April 4th, 2023
Scaled Cognition
- **Language, Structure, and Beyond** March 13th, 2023
Brandeis University – CS Department
- **Language, Structure, and Beyond** March 2nd, 2023
University of Cambridge – CS Department
- **Language, Structure, and Beyond** February 17th, 2023
George Mason University – George Mason NLP Group
- **Language, Structure, and Beyond** February 1st, 2023
Georgetown University – NERT Lab
- **Joint Universal Syntactic and Semantic Parsing** April 1st, 2022
Cornell University – Workshop on Meaning in Language

HONORS AND AWARDS

Honors and Awards

- First Class Honours in Cognitive Science 2018
- Dean's Honor List (top 10% of faculty) 2014–2015, 2016–2018
- Arts Undergraduate Research Internship Award 2016
- NSF Graduate Research Fellowship 2019

SERVICE

Reviewing

- **Area Chair:** ACL ARR (2024), EMNLP ARR (2024), NAACL ARR (2025), ACL ARR (2025), EMNLP ARR (2025), AACL ARR (2025)

- **Reviewer:** ACL (2021, 2022), NAACL 2022, EACL 2023, CoRL 2023, EMNLP 2023, NeurIPS (2024, 2025), WiNLP (2024, 2025), ICLR (2025, 2026), ICML 2025, ICCV 2025, TACL 2025
- **Program Committee:** UnImplicit 2022, SouthNLP 2024

Organizing

- **Workshop Organizer:** UnImplicit 2024
- **Conference Organizer:** EMNLP 2024 (Publicity Chair), ACL 2026 (Student Research Workshop Chair)
- **Session Chair:** NAACL 2025

PhD/Proposal/RPE Committees at UT Austin

- **PhD Proposals:** Anibal Solon Heinsfeld, Vanya Cohen, Liyan Tang
- **RPEs:** Fangcong Yin, Meng Chen

External Committees

- **PhD Proposals:** Amith Ananthram (Columbia)

Department Service

- **UT Austin Graduate Admissions Committee**, 2025-2026
- **UT Austin Linguistics Hiring Committee**, 2025-2026
- **CLSP Visit Weekend Committee**, 2018-2022
- **CLSP Admissions Committee**, 2022-2023
- **Application Mentoring Group**, 2022-2023

TEACHING

- **Instructor** at University of Texas at Austin
Communicative and Grounded AI Agents – CS395T Fall 2025
- **Guest Lecturer** at UNC Chapel Hill
Advanced Topics in Natural Language Processing (Aligning LLMs and Humans – COMP 790-158) Fall 2024
- **Guest Lecturer** at UNC Chapel Hill
Deep Learning (COMP 664) Fall 2024
- **Guest Lecturer** at UNC Chapel Hill
Computational Linguistics Fall 2023, Fall 2024
- **Teaching Assistant** at Johns Hopkins University
Artificial Intelligence (EN.601.464/664) Fall 2019

MENTORING

- Victor Wong
Undergraduate Honors Thesis (UT Austin) Summer 2025-Present
- Vaidehi Patil
Ph.D. Student (UNC) Fall 2024-Present
- Zaid Khan
Ph.D. Student (UNC) Fall 2024-Present
- Duy Nguyen
Fall 2024-Present

<i>Ph.D. Student (UNC)</i>	
• Amith Ananthram	Winter 2024-Present
<i>Ph.D. Student (Columbia)</i>	
• Ziyang Wang	Fall 2023-Fall 2024
<i>Ph.D. Student (UNC)</i>	
• Shoubin Yu	Fall 2023-Present
<i>Ph.D. Student (UNC)</i>	
• David Wan	Fall 2023-Present
<i>Ph.D. Student (UNC)</i>	
• Justin Chih-Yao Chen	Fall 2023-Present
<i>Ph.D. Student (UNC)</i>	
• Han Wang	Fall 2023-Present
<i>Ph.D. Student (UNC)</i>	
• Archiki Prasad	Summer 2023-Present
<i>Ph.D. Student (UNC)</i>	
• Victor Wang	Summer 2025-Present
<i>Undergrad (UT Austin)</i>	
Outstanding Undergraduate Researcher Award, CRA (Honorable Mention)	
• Tianyi Niu	Summer 2025-Present
<i>Undergrad (UNC)</i>	
• Hanqi Xiao	Summer 2024-Present
<i>Undergrad (UNC)</i>	
Outstanding Undergraduate Researcher Award, CRA (Honorable Mention)	
• Atin Pothviraj	Summer 2024-Present
<i>Undergrad (UNC)</i>	
• Nithin Sivakumaran	Summer 2024-Present
<i>Undergrad (UNC)</i>	
Outstanding Undergraduate Researcher Award, CRA (Honorable Mention)	
• Bryan Sukidi	Fall 2024-Fall 2025
<i>Undergrad (UNC)</i>	
• Evan Paces-Wiles	Fall 2024-Present
<i>Undergrad (UNC)</i>	
• Jio Choi	Fall 2024-Summer 2025
<i>Undergrad (UNC)</i>	
• Jimena Guallar-Blasco	Summer 2020-Summer 2024
<i>Undergrad/MS 2024 (JHU, Current: Charles River Analytics)</i>	
Outstanding Undergraduate Researcher Award, CRA (Honorable Mention)	
• Maryam Hashemzadeh	Summer 2023
<i>Ph.D. Student (MILA)</i>	
• Zhuohong (Zooey) He	Winter 2021-Spring 2021
<i>MSE, Spring 2021 (JHU, Current: Intuitive)</i>	
• Chenyu (Heidi) Zhang	Fall 2021-Fall 2022
<i>BS 2022, (Next: MS (Stanford), Current: Google Gemini)</i>	
• Shalaka Vaidya	Fall 2022
<i>MS 2023 (NYU)</i>	
• Yi Zhou	Winter 2022-Fall 2022

SKILLS

- **Programming (expert):** Python
- **Programming (proficient):** Bash, Java, Javascript
- **Libraries/Frameworks:** PyTorch, AllenNLP, Transformers, NLTK, numpy, MechanicalTurk, networkx, MXNet, React

LANGUAGES

- **Native:** English, German
- **Fluent:** French
- **Other:** Latin (reading/translation), Spanish (intermediate)