

Elias Stengel-Eskin

Website: esteng.github.io
Email: esteng@cs.utexas.edu
LinkedIn: [elias-stengel-eskin](https://www.linkedin.com/in/elias-stengel-eskin)
GitHub: github.com/esteng

EDUCATION

Johns Hopkins University

Baltimore, USA

Ph.D. in Computer Science – *Advisor: Benjamin Van Durme*

2018–2023

- Ph.D. thesis: “Modeling Meaning for Description and Interaction”
- supported by NSF Graduate Research Fellowship

Johns Hopkins University

Baltimore, USA

MSE in Computer Science – *Advisor: Benjamin Van Durme*

2018–2021

McGill University

Montréal, Canada

Bachelor of Arts and Sciences in Cognitive Science

2014–2018

- Minor: Linguistics
- First Class Honours
- Honours thesis: “Variational Bayesian Inference for Unsupervised Lexicon Discovery” – *Advisor: Timothy O’Donnell*

EXPERIENCE

The University of Texas at Austin

Assistant Professor, Department of Computer Science

August 2025–Present

Microsoft

Researcher (Consulting)

February 2025–Present

University of North Carolina, Chapel Hill

Postdoctoral Research Associate – *Advisor: Mohit Bansal*

August 2023–August 2025

Microsoft Research

PhD Research Intern – *Advisors: Marc-Alexandre Côté, Eric Yuan, Pierre-Yves Oudeyer*

March 2022–March 2023

Microsoft Research - Semantic Machines

PhD Research Intern – *Advisor: Yu Su*

May 2021 - August 2021

Montreal Computational and Quantitative Linguistics Lab

Research Assistant – *Advisor: Morgan Sonderegger*

2016–2018

PUBLICATIONS (PEER-REVIEWED)

- [1] Hanqi Xiao, Vaidehi Patil, Zaid Khan, Hyunji Lee, **Elias Stengel-Eskin**, and Mohit Bansal, “Playing along: Learning a double-agent defender for belief steering via theory of mind”, *Conference on Language Modeling*, 2026.
- [2] Joykirat Singh, Justin Chih-Yao Chen, Archiki Prasad, **Elias Stengel-Eskin**, Akshay Nambi, and Mohit Bansal, “Think right: Learning to mitigate under-over thinking via adaptive, attentive compression”, *Conference on Language Modeling*, 2026.

- [3] Justin Chih-Yao Chen, Sukwon Yun, **Elias Stengel-Eskin**, Tianlong Chen, and Mohit Bansal, “Symbolic Mixture-of-Experts: Adaptive skill-based routing for heterogeneous reasoning”, *The 43rd International Conference on Machine Learning (ICML)*, 2026.
- [4] David Wan, Han Wang, Ziyang Wang, **Elias Stengel-Eskin**, Hyunji Lee, and Mohit Bansal, “Multimodal fact-level attribution for verifiable reasoning”, *The 43rd International Conference on Machine Learning (ICML)*, 2026.
- [5] Hanqi Xiao, Vaidehi Patil, Hyunji Lee, **Elias Stengel-Eskin**, and Mohit Bansal, “Generalized correctness models: Learning calibrated and model-agnostic correctness predictors from historical patterns”, *The 43rd International Conference on Machine Learning (ICML)*, 2026.
- [6] Tianyi Niu, Justin Chih-Yao Chen, Genta Indra Winata, Shi-Xiong Zhang, Supriyo Chakraborty, Sambit Sahu, Yue Zhang, **Elias Stengel-Eskin**, and Mohit Bansal, “Routing with generated data: Annotation-free LLM skill estimation and expert selection”, *Proceedings of the Sixty-Fourth Annual Meeting of the Association for Computational Linguistics (ACL)*, 2026.
- [7] Jaewoo Lee, Archiki Prasad, Justin Chih-Yao Chen, Zaid Khan, **Elias Stengel-Eskin**, and Mohit Bansal, “PRInTS: Reward modeling for long-horizon information seeking”, *Proceedings of the Sixty-Fourth Annual Meeting of the Association for Computational Linguistics (ACL)*, 2026.
- [8] Duy Nguyen, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “GrAInS: Gradient-based attribution for inference-time steering of LLMs and VLMs”, *Proceedings of the Sixty-Fourth Annual Meeting of the Association for Computational Linguistics (ACL)*, 2026.
- [9] Nithin Sivakumaran, Justin Chih-Yao Chen, David Wan, Yue Zhang, Jaehong Yoon, **Elias Stengel-Eskin**, and Mohit Bansal, “DART: Leveraging multi-agent disagreement for tool recruitment in multimodal reasoning”, *Proceedings of the 2026 Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2026.
- [10] Joyce Horn Fonteles, Nithin Sivakumaran, Clayton Cohn, Austin Coursey, Shoubin Yu, **Elias Stengel-Eskin**, Ashwin TS, Mohit Bansal, and Gautam Biswas, “A novel approach to evaluating the effectiveness of large language models for multimodal analysis of embodied learning in classrooms”, in *Proceedings of the LAK26: 16th International Learning Analytics and Knowledge Conference*, 2026, pages 536–546.
- [11] Tianyi Niu, Jaemin Cho, **Elias Stengel-Eskin**, and Mohit Bansal, “RotBench: Evaluating multimodal large language models on identifying image rotation”, *Proceedings of the 2026 Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2026.
- [12] Hyunji Lee, Seunghyun Yoon, Yunjae Won, Hanseok Oh, Geewook Kim, Trung Bui, Franck Dernoncourt, **Elias Stengel-Eskin**, Mohit Bansal, and Minjoon Seo, “Context-informed grounding supervision”, *Proceedings of the 2026 Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2026.
- [13] Victor Wang and **Elias Stengel-Eskin**, “Calibrating verbalized confidence with self-generated distractors”, *The Fourteenth International Conference on Learning Representations (ICLR)*, 2026.
- [14] Hyunji Lee, Minseon Kim, Chinmay Singh, Matheus Pereira, Atharv Sonwane, Isadora White, **Elias Stengel-Eskin**, Mohit Bansal, Zhengyan Shi, Alessandro Sordoni, *et al.*, “Gistify! Codebase-level understanding via runtime execution”, *The Fourteenth International Conference on Learning Representations (ICLR)*, 2026.
- [15] Amith Ananthram, **Elias Stengel-Eskin**, Lorena A Bradford, Julia Demarest, Adam Purvis, Keith Krut, Robert Stein, Rina Elster Pantalony, Mohit Bansal, and Kathleen McKeown, “PoSh: Using scene graphs to guide LLMs-as-a-judge for detailed image descriptions”, *The Fourteenth International Conference on Learning Representations (ICLR)*, 2026.

- [16] Zaid Khan, Archiki Prasad, **Elias Stengel-Eskin**, Jaemin Cho, and Mohit Bansal, “One life to learn: Inferring symbolic world models for stochastic environments from unguided exploration”, *The Fourteenth International Conference on Learning Representations (ICLR)*, 2026.
- [17] Jio Choi, Mohit Bansal, and **Elias Stengel-Eskin**, “Language models identify ambiguities and exploit loopholes”, *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2025.
- [18] Justin Chih-Yao Chen, Archiki Prasad, Swarnadeep Saha, **Elias Stengel-Eskin**, and Mohit Bansal, “Magicore: Multi-agent, iterative, coarse-to-fine refinement for reasoning”, *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2025.
- [19] Duy Nguyen, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “Laser: Learning to adaptively select reward models with multi-armed bandits”, *Advances in Neural Information Processing Systems* 39, 2025.
- [20] Eduardo Calò, Guanyi Chen, **Elias Stengel-Eskin**, Albert Gatt, and Kees van Deemter, “Incorporating formulaicness in the automatic evaluation of naturalness: A case study in logic-to-text generation”, in *Proceedings of the 18th International Natural Language Generation Conference*, 2025, pages 352–365.
- [21] Nithin Sivakumaran, Chia-Yu Yang, Abhay Zala, Shoubin Yu, Daeun Hong, Xiaotian Zou, **Elias Stengel-Eskin**, Dan Carpenter, Wookhee Min, Cindy Hmelo-Silver, *et al.*, “A multimodal classroom video question-answering framework for automated understanding of collaborative learning”, in *Proceedings of the 27th International Conference on Multimodal Interaction*, 2025, pages 516–525.
- [22] Han Wang, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “Retrieval-augmented generation with conflicting evidence”, *Conference on Language Modeling (COLM)*, 2025.
- [23] Hanqi Xiao, Yi-Lin Sung, **Elias Stengel-Eskin**, and Mohit Bansal, “Task-circuit quantization: Leveraging knowledge localization and interpretability for compression”, *Conference on Language Modeling (COLM)*, 2025.
- [24] Archiki Prasad, **Elias Stengel-Eskin**, Justin Chih-Yao Chen, Zaid Khan, and Mohit Bansal, “Learning to generate unit tests for automated debugging”, *Conference on Language Modeling (COLM)*, 2025.
- [25] David Wan, Eran Hirsch, **Elias Stengel-Eskin**, Ido Dagan, and Mohit Bansal, “Generationprograms: Fine-grained attribution with executable programs”, *Conference on Language Modeling (COLM)*, 2025.
- [26] Atin Pothiraj, **Elias Stengel-Eskin**, Jaemin Cho, and Mohit Bansal, “Capture: Evaluating spatial reasoning in vision language models via occluded object counting”, *Proceedings of the International Conference on Computer Vision, (ICCV)*, 2025.
- [27] Duy Nguyen, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “Multi-attribute steering of language models via targeted intervention”, *Proceedings of the Sixty-Third Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025.
- [28] Eran Hirsch, Aviv Slobodkin, David Wan, **Elias Stengel-Eskin**, Mohit Bansal, and Ido Dagan, “Laquer: Localized attribution queries in content-grounded generation”, *Proceedings of the Sixty-Third Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025.
- [29] Ziyang Wang*, Shoubin Yu*, **Elias Stengel-Eskin***, Jaehong Yoon, Feng Cheng, Gedas Bertasius, and Mohit Bansal, “VideoTree: Adaptive tree-based video representation for LLM reasoning on long videos”, *Proceedings of the Forty-Second Annual IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
- [30] Zaid Khan, **Elias Stengel-Eskin**, Jaemin Cho, and Mohit Bansal, “DataEnvGym: Data generation agents in teacher environments with student feedback”, *The Thirteenth International Conference on Learning Representations (ICLR) Spotlight*, 2025.

- [31] Swarnadeep Saha, Archiki Prasad, Justin Chih-Yao Chen, Peter Hase, **Elias Stengel-Eskin**, and Mohit Bansal, “System-1.x: Learning to balance fast and slow planning with language models”, *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- [32] Amith Ananthram, **Elias Stengel-Eskin**, Carl Vondrick, Mohit Bansal, and Kathleen McKeown, “See it from my perspective: Diagnosing the western cultural bias of large vision-language models in image understanding”, *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- [33] **Elias Stengel-Eskin**, Peter Hase, and Mohit Bansal, “Teaching models to balance resisting and accepting persuasion”, *Proceedings of the Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)*, 2025.
- [34] David Wan, Justin Chih-Yao Chen, **Elias Stengel-Eskin**, and Mohit Bansal, “MAMM-Refine: A recipe for improving faithfulness in generation with multi-agent collaboration”, *Proceedings of the Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)*, 2025.
- [35] Han Wang, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “AdaCAD: Adaptively decoding to balance conflicts between contextual and parametric knowledge”, *Proceedings of the Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)*, 2025.
- [36] **Elias Stengel-Eskin**, Peter Hase, and Mohit Bansal, “LACIE: Listener-aware finetuning for confidence calibration in large language models”, *Advances in Neural Information Processing Systems 38*, 2024.
- [37] Jinhao Duan, Renming Zhang, James Diffenderfer, Bhavya Kailkhura, Lichao Sun, **Elias Stengel-Eskin**, Mohit Bansal, Tianlong Chen, and Kaidi Xu, “GTBench: Uncovering the strategic reasoning limitations of LLMs via game-theoretic evaluations”, *Advances in Neural Information Processing Systems 38*, 2024.
- [38] Ansel Blume, Khanh Duy Nguyen, Zhenhailong Wang, Yangyi Chen, Michal Shlapentokh-Rothman, Xiaomeng Jin, Jeonghwan Kim, ..., **Stengel-Eskin-Elias**, *et al.*, “MIRACLE: An online, explainable multimodal interactive concept learning system”, in *Proceedings of the 32nd ACM International Conference on Multimedia*, 2024.
- [39] Peter Hase, Thomas Hofweber, Xiang Zhou, **Elias Stengel-Eskin**, and Mohit Bansal, “Fundamental problems with model editing: How should rational belief revision work in LLMs?”, *Transactions on Machine Learning*, 2024.
- [40] **Elias Stengel-Eskin**, Archiki Prasad, and Mohit Bansal, “ReGAL: Refactoring programs to discover generalizable abstractions”, *The 41st International Conference on Machine Learning (ICML)*, 2024.
- [41] Justin Chih-Yao Chen, Swarnadeep Saha, **Elias Stengel-Eskin**, and Mohit Bansal, “MAGDi: Structured distillation of multi-agent interaction graphs improves reasoning in smaller language models”, *The 41st International Conference on Machine Learning (ICML)*, 2024.
- [42] Haotian Fu, Pratyusha Sharma, **Elias Stengel-Eskin**, George Konidaris, Nicolas Le Roux, Marc-Alexandre Côté, and Xingdi Yuan, “Language-guided skill learning with temporal variational inference”, *The 41st International Conference on Machine Learning (ICML)*, 2024.
- [43] David Wan, Jaemin Cho, **Elias Stengel-Eskin**, and Mohit Bansal, “Contrastive region guidance: Improving grounding in vision-language models without training”, *The 18th European Conference on Computer Vision (ECCV)*, 2024.
- [44] Han Wang*, Archiki Prasad*, **Elias Stengel-Eskin***, and Mohit Bansal, “Soft self-consistency improves language model agents”, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- [45] Maryam Hashemzadeh, **Elias Stengel-Eskin**, Sarath Chandar, and Marc-Alexandre Cote, “Sub-goal distillation: A method to improve small language agents”, *The 3rd Conference on Lifelong Learning Agents (CoLLAs)*, 2024.

- [46] **Elias Stengel-Eskin**, Kyle Rawlins, and Benjamin Van Durme, “Zero and few-shot semantic parsing with ambiguous inputs”, in *The Twelfth International Conference on Learning Representations (ICLR)*, 2024.
- [47] Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “Rephrase, augment, reason: Visual grounding of questions for vision-language models”, in *The Twelfth International Conference on Learning Representations (ICLR)*, 2024.
- [48] **Elias Stengel-Eskin** and Benjamin Van Durme, “Did You Mean...? Confidence-based Trade-offs in Semantic Parsing”, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [49] **Elias Stengel-Eskin** and Benjamin Van Durme, “Calibrated Interpretation: Confidence Estimation in Semantic parsing”, *Transactions of the Association for Computational Linguistics (TACL)*, 2023.
- [50] **Elias Stengel-Eskin**, Jimena Guallar-Blasco, Yi Zhou, and Benjamin Van Durme, “Why Did the Chicken Cross the Road? Rephrasing and Analyzing Ambiguous Questions in VQA”, *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023.
- [51] Zhuowan Li, Xingrui Wang, **Elias Stengel-Eskin**, Adam Kortylewski, Wufei Ma, Benjamin Van Durme, and Alan Yuille, “Super-CLEVR: A Virtual Benchmark to Diagnose Domain Robustness in Visual Reasoning”, *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [52] **Elias Stengel-Eskin** and Benjamin Van Durme, “The Curious Case of Control”, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022.
- [53] **Elias Stengel-Eskin**, Emmanouil Antonios Platanios, Adam Pauls, Sam Thomson, Hao Fang, Benjamin Van Durme, Jason Eisner, and Yu Su, “When More Data Hurts: A Troubling Quirk in Developing Broad-Coverage Natural Language Understanding Systems”, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022.
- [54] Chenyu Zhang, Benjamin Van Durme, Zhuowan Li, and **Elias Stengel-Eskin**, “Visual Commonsense in Pretrained Unimodal and Multimodal Models”, in *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, Seattle, Washington: Association for Computational Linguistics, 2022.
- [55] **Elias Stengel-Eskin**, Andrew Hundt, Zhuohong He, Aditya Murali, Nakul Gopalan, Matthew Gombolay, and Gregory D. Hager, “Guiding Multi-Step Rearrangement Tasks with Natural Language Instructions”, in *5th Annual Conference on Robot Learning (CoRL)*, 2021.
- [56] Zhuowan Li, **Elias Stengel-Eskin**, Yixiao Zhang, Cihang Xie, Quan Tran, Benjamin Van Durme, and Alan Yuille, “Calibrating Concepts and Operations: Towards Symbolic Reasoning on Real Images”, in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021.
- [57] **Elias Stengel-Eskin**, Jimena Guallar-Blasco, and Benjamin Van Durme, “Human-Model Divergence in the Handling of Vagueness”, in *Proceedings of the 1st Workshop on Understanding Implicit and Underspecified Language (UnImplicit)*, Online: Association for Computational Linguistics, 2021.
- [58] **Elias Stengel-Eskin**, Kenton Murray, Sheng Zhang, Aaron Steven White, and Benjamin Van Durme, “Joint Universal Syntactic and Semantic Parsing”, *Transactions of the Association for Computational Linguistics (TACL)*, 2021.
- [59] **Elias Stengel-Eskin**, Jimena Guallar-Blasco, and Benjamin Van Durme, “Exploring Human-Model Divergence Through Vagueness”, *Proceedings of the Society for Computation in Linguistics (SCiL)*, 2021, ***Abstract**.
- [60] Ryan Culkin, J. Edward Hu, **Elias Stengel-Eskin**, Guanghui Qin, and Benjamin Van Durme, “Iterative Paraphrastic Augmentation with Discriminative Span Alignment”, *Transactions of the Association for Computational Linguistics (TACL)*, volume 9, 2021.

- [61] **Elias Stengel-Eskin**, Aaron Steven White, Sheng Zhang, and Benjamin Van Durme, “Universal Decompositional Semantic Parsing”, in *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2020.
- [62] Aaron Steven White, **Elias Stengel-Eskin**, Siddharth Vashishtha, Venkata Subrahmanyam Govindarajan, Dee Ann Reisinger, Tim Vieira, Keisuke Sakaguchi, Sheng Zhang, Francis Ferraro, Rachel Rudinger, *et al.*, “The Universal Decompositional Semantics Dataset and Decomp Toolkit”, in *Proceedings of The 12th Language Resources and Evaluation Conference*, 2020.
- [63] **Elias Stengel-Eskin**, Tzu-Ray Su, Matt Post, and Benjamin Van Durme, “A Discriminative Neural Model for Cross-Lingual Word Alignment”, in *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, 2019.
- [64] Michael McAuliffe, **Elias Stengel-Eskin**, Michaela Socolof, and Morgan Sonderegger, “Polyglot and Speech Corpus Tools: A System for Representing, Integrating, and Querying Speech Corpora.” In *INTERSPEECH*, 2017.

PUBLICATIONS (OTHER)

1. Hyunji Lee, Justin Chih-Yao Chen, Joykirat Singh, Zaid Khan, **Elias Stengel-Eskin**, Mohit Bansal, “LongMINT: Evaluating Memory under Multi-Target Interference in Long-Horizon Agent Systems” (2026), *Under Review*
2. Duy Nguyen, Hanqi Xiao, Archiki Prasad, Zaid Khan, Anirban Das, Austin Zhang, Sambit Sahu, Hyunji Lee, **Elias Stengel-Eskin**, Mohit Bansal, “AVSD: Adaptive-View Self-Distillation by Balancing Consensus and Teacher-Specific Privileged Signals” (2026), *Under Review*
3. Joykirat Singh, Zaid Khan, Archiki Prasad, Justin Chih-Yao Chen, Akshay Nambi, Hyunji Lee, **Elias Stengel-Eskin**, Mohit Bansal, “Agent-BRACE: Decoupling Beliefs from Actions in Long-Horizon Tasks via Verbalized State Uncertainty” (2026) *Under Review*
4. Saeid Asgari Taghanaki, Rakshanda Agarwal, Bruce Sun, Rohan Jha, **Elias Stengel-Eskin**, Sara Malvar, Rui Ying, Yifei Xu, Guilherme Potje, Tusher Chakraborty, Leonardo de Oliveira Nunes, Ranveer Chandra, Emre Kiciman, “Diagnosing Capability Gaps in Fine-Tuning Data” (2026) *Under Review*
5. Han Wang, David Wan, Hyunji Lee, Thinh Pham, Mikaela Cankosyan, Weiyuan Chen, **Elias Stengel-Eskin**, Tu Vu, Mohit Bansal, “MERRIN: A Benchmark for Multimodal Evidence Retrieval and Reasoning in Noisy Web Environments” (2026), *Under Review*
6. Justin Chih-Yao Chen, Archiki Prasad, Zaid Khan, Joykirat Singh, Runchu Tian, **Elias Stengel-Eskin**, Mohit Bansal, “Cog-DRIFT: Exploration on Adaptively Reformulated Instances Enables Learning from Hard Reasoning Problems” (2026), *Under Review*
7. Marvin Limpijankit, Milad Alshomary, Yassin Oulad Daoud, Amith Ananthram, Tim Trombley, **Elias Stengel-Eskin**, Mohit Bansal, Noam M. Elcott, Kathleen McKeown, “Does AI See like Art Historians? Interpreting How Vision Language Models Recognize Artistic Style” (2026), *Under Review*
8. Vaidehi Patil, **Elias Stengel-Eskin**, Mohit Bansal, “The Sum Leaks More Than Its Parts: Compositional Privacy Risks and Mitigations in Multi-Agent Collaboration” (2025), *Under Review*
9. Nithin Sivakumaran, Shoubin Yu, Hyunji Lee, Yue Zhang, Ali Payani, Mohit Bansal, **Elias Stengel-Eskin**, “Balancing Faithfulness and Performance in Reasoning via Multi-Listener Soft Execution” (2026), *Under Review*
10. Duy Nguyen, Hanqi Xiao, Archiki Prasad, **Elias Stengel-Eskin**, Hyunji Lee, Mohit Bansal, “Conflict-Resolving and Sharpness-Aware Minimization for Generalized Knowledge Editing with Multiple Updates” (2026), *Under Review*
11. Haotian Zhai, **Elias Stengel-Eskin**, Pratik Patil, Liu Leqi, “Evaluating Stochasticity in Deep Research Agents” (2026), *Under Review*

12. David Wan, Han Wang, **Elias Stengel-Eskin**, Jaemin Cho, Mohit Bansal, “CLaMR: Contextualized Late-Interaction for Multimodal Content Retrieval” (2025), *Under Review*
13. Emmanouil Zaranis, António Farinhas, Saul Santos, Beatriz Canaverde, Miguel Moura Ramos, Aditya K Surikuchi, André Viveiros, Baohao Liao, Elena Bueno-Benito,..., Nithin Sivakumaran, **Elias Stengel-Eskin**, et al., “Movie Facts and Fibs (MF²): A Benchmark for Long Movie Understanding” (2025), *Under Review*
14. Ron Eliav, Arie Cattan, Eran Hirsch, Shahaf Bassan, **Elias Stengel-Eskin**, Mohit Bansal, Ido Dagan, “CLATTER: Comprehensive Entailment Reasoning for Hallucination Detection” (2025), *Under Review*
15. Zaid Khan, **Elias Stengel-Eskin**, Archiki Prasad, Jaemin Cho, Mohit Bansal, “Executable Functional Abstractions: Inferring Generative Programs for Advanced Math Problems” (2025), *Under Review*
16. Vaidehi Patil, **Elias Stengel-Eskin**, Mohit Bansal, “UPCORE: Utility-Preserving Coreset Selection for Balanced Unlearning” (2025), *Under Review*
17. Yue Huang, Chujie Gao, Siyuan Wu, ..., **Elias Stengel-Eskin**, ..., Bo Li, Xiangliang Zhang, “On the Trustworthiness of Generative Foundation Models: Guideline, Assessment, and Perspective” (2025), *Preprint*
18. Thomas Hofweber, Peter Hase, **Elias Stengel-Eskin**, Mohit Bansal, “Are language models rational? The case of coherence norms and belief revision” (2024), *Preprint*
19. **Elias Stengel-Eskin**, “Modeling Meaning for Description and Interaction”, (2023) **Ph.D. Thesis**
20. Jimena Guallar-Blasco, **Elias Stengel-Eskin**, Benjamin Van Durme, “Analyzing Question Ambiguity in *Why*-Questions” (2023), *Mid-Atlantic Student Colloquium on Speech, Language and Learning (MASC-SLL 2023)* ***Abstract**
21. Shalaka Vaidya, **Elias Stengel-Eskin**, Joao Sedoc, “Automatic Evaluation of Chit-chat via Semantic Parsing” (2022), *Mid-Atlantic Student Colloquium on Speech, Language and Learning (MASC-SLL 2022)* ***Abstract**
22. Yunmo Chen, Seth Ebner, Tongfei Chen, Patrick Xia, **Elias Stengel-Eskin**, Tzu-Ray Su, J. Edward Hu, Nils Holzenberger, Ryan Culkin, Craig Harman, Max Thomas, Thomas Lippincott, Aaron Steven White, Kyle Rawlins, Benjamin Van Durme, “NIST TAC SM-KBP 2019 System Description: JHU/UR Framework”, (2019)
23. **Elias Stengel-Eskin**, “Variational Bayesian Inference for Unsupervised Lexicon Discovery” (2017), **Undergraduate Honours Thesis**

FUNDING

Received

- **IARPA BENGAL** 2025-2027
\$290k. Role: Postdoc
- **NSF NAIRR Pilot Grant** 2024-2025
Supervising Calibration for Safe and Trustworthy Artificial Intelligence Systems
2,500 GPU Hours. Role: **PI**
- **Cisco Gift** 2025
Improving LLM Reliability and Trustworthiness via Overconfidence Reduction and Conversation
\$75K. Role: Co-PI
- **CapitalOne Gift** 2025
Multi-Agent Systems and Mixtures-of-Experts for Robust Generation
Role: Co-PI
- **Center for AI Safety Compute Grant** 2024
Listener-Aware Confidence Calibration for Improving LLM Reliability and Trustworthiness
Unlimited A100 GPU access. Role: Co-PI
- **NSF Graduate Research Fellowship** 2018-2023
Role: Student

NEWS COVERAGE

1. *The importance of ambiguity*, Johns Hopkins University CS News, 04/15/2024
2. *Speaking easy: A simple method to clarify voice commands*, Johns Hopkins University CS News, 01/08/2024

INVITED TALKS/PANELS

- **Multi-Model Training for Multi-Agent Communication Skills** 2026
Keynote – ICLR 2026 Workshop on Multi-Agent Learning in Generative AI
- **Multi-Model Training for Multi-Agent Communication Skills** 2026
Harrington Annual Symposium
- **Multi-Model Training for Multi-Agent Communication Skills** 2026
Stamina Working Group
- **Multi-Model Training for Multi-Agent Communication Skills** 2026
Johns Hopkins University, DSAI Seminar
- **Multi-Model Training for Multi-Agent Communication Skills** 2026
University of Pennsylvania, Clunch
- **Multi-Model Training for Multi-Agent Communication Skills** March 13, 2026
Workshop on Large Language Models, Rice University
- **Panel: Reimagining Work in the Age of AI Agents** March 4, 2026
Texas Symposium on Machine Learning, Responsible AI & Robotics
- **Towards Granular, Robust, Calibrated Judgments via Skill-Based and Persuasion-Balanced Reasoning** December 2, 2025
NeurIPS Expo Workshop on Exploring Trust and Reliability in LLM Evaluation
- **Panel: Agentic AI in Wireless Networking and Communications** November 14, 2025
UT Austin Wireless Networking and Communications Group
- **Panel: The true impact of AI on the labor market** September 30, 2025
Hire EZ
- **Communication, Calibration, and Grounding for Collaborative AI Agents** January 8, 2025
Duke University
- **Learning Pragmatic Skills through Multi-Agent Training** 2024
UNC Chapel Hill – Linguistics Department
- **Confidence-based Rephrasing, Refinement, and Selection** March 22nd, 2024
Keynote – UncertainNLP Workshop, EACL 2024
- **Ambiguity in NLP** October 26th, 2023
UNC Chapel Hill – Linguistics Department
- **Uncertainty and Ambiguity in Semantic Structures** October 6th, 2023
Emory University
- **Language, Structure, and Beyond** May 9th, 2023
UNC Chapel Hill – MURGe Group
- **Language, Structure, and Beyond** April 4th, 2023
Scaled Cognition
- **Language, Structure, and Beyond** March 13th, 2023
Brandeis University – CS Department
- **Language, Structure, and Beyond** March 2nd, 2023
University of Cambridge – CS Department

- **Language, Structure, and Beyond** February 17th, 2023
George Mason University – George Mason NLP Group
- **Language, Structure, and Beyond** February 1st, 2023
Georgetown University – NERT Lab
- **Joint Universal Syntactic and Semantic Parsing** April 1st, 2022
Cornell University – Workshop on Meaning in Language

HONORS AND AWARDS

Honors and Awards

- First Class Honours in Cognitive Science 2018
- Dean's Honor List (top 10% of faculty) 2014–2015, 2016–2018
- Arts Undergraduate Research Internship Award 2016
- NSF Graduate Research Fellowship 2019

SERVICE

Reviewing

- **Area Chair:** ARR (2024: 2x, 2025: 4x, 2026: 2x), COLM (2026)
- **Reviewer:** ACL (2021, 2022), NAACL 2022, EACL 2023, CoRL 2023, EMNLP 2023, NeurIPS (2024, 2025), WinLP (2024, 2025), ICLR (2025, 2026), ICML (2025, 2026 – Gold Reviewer), ICCV 2025, TACL (2025, 2026), CVPR 2026, ECCV 2026
- **Program Committee:** UnImplicit 2022, SouthNLP 2024
- **Senior Area Chair:** EMNLP 2026

Organizing

- **Workshop Organizer:** UnImplicit 2024
- **Conference Organizer:** EMNLP 2024 (Publicity Chair), ACL 2026 (Student Research Workshop Chair)
- **Session Chair:** NAACL 2025

PhD/Proposal/RPE Committees at UT Austin

- **Thesis Committees:** Liyan Tang
- **PhD Proposals:** Anibal Solon Heinsfeld, Vanya Cohen, Liyan Tang
- **RPEs:** Fangcong Yin, Meng Chen, Parikshit Bansal

External Committees

- **PhD Proposals:** Amith Ananthram (Columbia)

Department Service

- **UT Austin Graduate Admissions Committee,** 2025-2026
- **UT Austin Linguistics Hiring Committee,** 2025-2026
- **CLSP Visit Weekend Committee,** 2018-2022
- **CLSP Admissions Committee,** 2022-2023
- **Application Mentoring Group,** 2022-2023

TEACHING

- **Instructor** at University of Texas at Austin Fall 2025
Communicative and Grounded AI Agents – CS395T
- **Guest Lecturer** at UNC Chapel Hill Fall 2024
Advanced Topics in Natural Language Processing (Aligning LLMs and Humans – COMP 790-158)
- **Guest Lecturer** at UNC Chapel Hill Fall 2024
Deep Learning (COMP 664)
- **Guest Lecturer** at UNC Chapel Hill Fall 2023, Fall 2024
Computational Linguistics
- **Teaching Assistant** at Johns Hopkins University Fall 2019
Artificial Intelligence (EN.601.464/664)

MENTORING

- Vaidehi Patil Fall 2024-Present
Ph.D. Student (UNC)
- Zaid Khan Fall 2024-Present
Ph.D. Student (UNC)
- Duy Nguyen Fall 2024-Present
Ph.D. Student (UNC)
- Amith Ananthram Winter 2024-Present
Ph.D. Student (Columbia)
- Ziyang Wang Fall 2023-Fall 2024
Ph.D. Student (UNC)
- Shoubin Yu Fall 2023-Present
Ph.D. Student (UNC)
- David Wan Fall 2023-Present
Ph.D. Student (UNC)
- Justin Chih-Yao Chen Fall 2023-Present
Ph.D. Student (UNC)
- Han Wang Fall 2023-Present
Ph.D. Student (UNC)
- Archiki Prasad Summer 2023-Present
Ph.D. Student (UNC)
- Victor Wang Summer 2025-Present
Undergrad (UT Austin)
Outstanding Undergraduate Researcher Award, CRA (Honorable Mention), NSF GRFP
- Tianyi Niu Summer 2025-Present
MS (UNC)
- Hanqi Xiao Summer 2024-Present
Undergrad (UNC, Current: NYU PhD student)
Outstanding Undergraduate Researcher Award, CRA (Honorable Mention), NSF GRFP
- Atin Pothviraj Summer 2024-Present
Undergrad (UNC)
- Nithin Sivakumaran Summer 2024-Present

Undergrad (UNC, Current: UMD PhD student)

Outstanding Undergraduate Researcher Award, CRA (Honorable Mention), NSF GRFP

- Bryan Sukidi Fall 2024-Fall 2025
Undergrad (UNC)
- Evan Paces-Wiles Fall 2024-Present
Undergrad (UNC)
- Jio Choi Fall 2024-Summer 2025
Undergrad (UNC)
- Jimena Guallar-Blasco Summer 2020-Summer 2024
Undergrad/MS 2024 (JHU, Current: Charles River Analytics)
Outstanding Undergraduate Researcher Award, CRA (Honorable Mention)
- Maryam Hashemzadeh Summer 2023
Ph.D. Student (MILA)
- Zhuohong (Zooey) He Winter 2021-Spring 2021
MSE, Spring 2021 (JHU, Current: Intuitive)
- Chenyu (Heidi) Zhang Fall 2021-Fall 2022
BS 2022, (JHU, Next: MS (Stanford), Current: Google Gemini)
- Shalaka Vaidya Fall 2022
MS 2023 (NYU)
- Yi Zhou Winter 2022-Fall 2022
MS 2022 (JHU)
- Vicky Zeng Winter 2021-Winter 2022
Ph.D. Student (JHU)

SKILLS

- **Programming (expert):** Python
- **Programming (proficient):** Bash, Java, Javascript
- **Libraries/Frameworks:** PyTorch, AllenNLP, Transformers, NLTK, numpy, MechanicalTurk, networkx, MXNet, React

LANGUAGES

- **Native:** English, German
- **Fluent:** French
- **Other:** Latin (reading/translation), Spanish (intermediate)