

Elias Stengel-Eskin

Website: esteng.github.io
Email: esteng@cs.utexas.edu
LinkedIn: [elias-stengel-eskin](https://www.linkedin.com/in/elias-stengel-eskin)
GitHub: github.com/esteng

EDUCATION

Johns Hopkins University

Ph.D. in Computer Science – *Advisor: Benjamin Van Durme*

Baltimore, USA

2018–2023

- Ph.D. thesis: “Modeling Meaning for Description and Interaction”
- supported by NSF Graduate Research Fellowship

Johns Hopkins University

MSE in Computer Science – *Advisor: Benjamin Van Durme*

Baltimore, USA

2018–2021

McGill University

Bachelor of Arts and Sciences in Cognitive Science

Montréal, Canada

2014–2018

- Minor: Linguistics
- First Class Honours
- Honours thesis: “Variational Bayesian Inference for Unsupervised Lexicon Discovery” – *Advisor: Timothy O’Donnell*

EXPERIENCE

The University of Texas at Austin

Assistant Professor, Department of Computer Science

August 2025–Present

Microsoft

Researcher (Consulting)

February 2025–Present

University of North Carolina, Chapel Hill

Postdoctoral Research Associate – *Advisor: Mohit Bansal*

August 2023–August 2025

Microsoft Research

PhD Research Intern – *Advisors: Marc-Alexandre Côté, Eric Yuan, Pierre-Yves Oudeyer*

March 2022–March 2023

Microsoft Research - Semantic Machines

PhD Research Intern – *Advisor: Yu Su*

May 2021 - August 2021

Montreal Computational and Quantitative Linguistics Lab

Research Assistant – *Advisor: Morgan Sonderegger*

2016–2018

PUBLICATIONS (PEER-REVIEWED)

- [1] Tianyi Niu, Justin Chih-Yao Chen, Genta Indra Winata, Shi-Xiong Zhang, Supriyo Chakraborty, Sambit Sahu, Yue Zhang, **Elias Stengel-Eskin**, and Mohit Bansal, “Routing with generated data: Annotation-free LLM skill estimation and expert selection”, *Proceedings of the Sixty-Fourth Annual Meeting of the Association for Computational Linguistics (ACL)*, 2026.
- [2] Jaewoo Lee, Archiki Prasad, Justin Chih-Yao Chen, Zaid Khan, **Elias Stengel-Eskin**, and Mohit Bansal, “PRInTS: Reward modeling for long-horizon information seeking”, *Proceedings of the Sixty-Fourth Annual Meeting of the Association for Computational Linguistics (ACL)*, 2026.

- [3] Duy Nguyen, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “GrAInS: Gradient-based attribution for inference-time steering of LLMs and VLMs”, *Proceedings of the Sixty-Fourth Annual Meeting of the Association for Computational Linguistics (ACL)*, 2026.
- [4] Nithin Sivakumaran, Justin Chih-Yao Chen, David Wan, Yue Zhang, Jaehong Yoon, **Elias Stengel-Eskin**, and Mohit Bansal, “DART: Leveraging multi-agent disagreement for tool recruitment in multimodal reasoning”, *Proceedings of the 2026 Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2026.
- [5] Tianyi Niu, Jaemin Cho, **Elias Stengel-Eskin**, and Mohit Bansal, “RotBench: Evaluating multimodal large language models on identifying image rotation”, *Proceedings of the 2026 Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2026.
- [6] Hyunji Lee, Seunghyun Yoon, Yunjae Won, Hanseok Oh, Geewook Kim, Trung Bui, Franck Dernoncourt, **Elias Stengel-Eskin**, Mohit Bansal, and Minjoon Seo, “Context-informed grounding supervision”, *Proceedings of the 2026 Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2026.
- [7] Victor Wang and **Elias Stengel-Eskin**, “Calibrating verbalized confidence with self-generated distractors”, *The Fourteenth International Conference on Learning Representations (ICLR)*, 2026.
- [8] Hyunji Lee, Minseon Kim, Chinmay Singh, Matheus Pereira, Atharv Sonwane, Isadora White, **Elias Stengel-Eskin**, Mohit Bansal, Zhengyan Shi, Alessandro Sordoni, *et al.*, “Gistify! Codebase-level understanding via runtime execution”, *The Fourteenth International Conference on Learning Representations (ICLR)*, 2026.
- [9] Amith Ananthram, **Elias Stengel-Eskin**, Lorena A Bradford, Julia Demarest, Adam Purvis, Keith Krut, Robert Stein, Rina Elster Pantalony, Mohit Bansal, and Kathleen McKeown, “PoSh: Using scene graphs to guide LLMs-as-a-judge for detailed image descriptions”, *The Fourteenth International Conference on Learning Representations (ICLR)*, 2026.
- [10] Zaid Khan, Archiki Prasad, **Elias Stengel-Eskin**, Jaemin Cho, and Mohit Bansal, “One life to learn: Inferring symbolic world models for stochastic environments from unguided exploration”, *The Fourteenth International Conference on Learning Representations (ICLR)*, 2026.
- [11] Jio Choi, Mohit Bansal, and **Elias Stengel-Eskin**, “Language models identify ambiguities and exploit loopholes”, *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2025.
- [12] Justin Chih-Yao Chen, Archiki Prasad, Swarnadeep Saha, **Elias Stengel-Eskin**, and Mohit Bansal, “Magicore: Multi-agent, iterative, coarse-to-fine refinement for reasoning”, *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2025.
- [13] Duy Nguyen, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “Laser: Learning to adaptively select reward models with multi-armed bandits”, *Advances in Neural Information Processing Systems* 39, 2025.
- [14] Eduardo Calò, Guanyi Chen, **Elias Stengel-Eskin**, Albert Gatt, and Kees van Deemter, “Incorporating formulaicness in the automatic evaluation of naturalness: A case study in logic-to-text generation”, in *Proceedings of the 18th International Natural Language Generation Conference*, 2025, pages 352–365.
- [15] Nithin Sivakumaran, Chia-Yu Yang, Abhay Zala, Shoubin Yu, Daeun Hong, Xiaotian Zou, **Elias Stengel-Eskin**, Dan Carpenter, Wookhee Min, Cindy Hmelo-Silver, *et al.*, “A multimodal classroom video question-answering framework for automated understanding of collaborative learning”, in *Proceedings of the 27th International Conference on Multimodal Interaction*, 2025, pages 516–525.
- [16] Han Wang, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “Retrieval-augmented generation with conflicting evidence”, *Conference on Language Modeling (COLM)*, 2025.

- [17] Hanqi Xiao, Yi-Lin Sung, **Elias Stengel-Eskin**, and Mohit Bansal, “Task-circuit quantization: Leveraging knowledge localization and interpretability for compression”, *Conference on Language Modeling (COLM)*, 2025.
- [18] Archiki Prasad, **Elias Stengel-Eskin**, Justin Chih-Yao Chen, Zaid Khan, and Mohit Bansal, “Learning to generate unit tests for automated debugging”, *Conference on Language Modeling (COLM)*, 2025.
- [19] David Wan, Eran Hirsch, **Elias Stengel-Eskin**, Ido Dagan, and Mohit Bansal, “Generationprograms: Fine-grained attribution with executable programs”, *Conference on Language Modeling (COLM)*, 2025.
- [20] Atin Pothiraj, **Elias Stengel-Eskin**, Jaemin Cho, and Mohit Bansal, “Capture: Evaluating spatial reasoning in vision language models via occluded object counting”, *Proceedings of the International Conference on Computer Vision, (ICCV)*, 2025.
- [21] Duy Nguyen, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “Multi-attribute steering of language models via targeted intervention”, *Proceedings of the Sixty-Third Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025.
- [22] Eran Hirsch, Aviv Slobodkin, David Wan, **Elias Stengel-Eskin**, Mohit Bansal, and Ido Dagan, “Laquer: Localized attribution queries in content-grounded generation”, *Proceedings of the Sixty-Third Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025.
- [23] Ziyang Wang*, Shoubin Yu*, **Elias Stengel-Eskin***, Jaehong Yoon, Feng Cheng, Gedas Bertasius, and Mohit Bansal, “VideoTree: Adaptive tree-based video representation for LLM reasoning on long videos”, *Proceedings of the Forty-Second Annual IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
- [24] Zaid Khan, **Elias Stengel-Eskin**, Jaemin Cho, and Mohit Bansal, “DataEnvGym: Data generation agents in teacher environments with student feedback”, *The Thirteenth International Conference on Learning Representations (ICLR) Spotlight*, 2025.
- [25] Swarnadeep Saha, Archiki Prasad, Justin Chih-Yao Chen, Peter Hase, **Elias Stengel-Eskin**, and Mohit Bansal, “System-1.x: Learning to balance fast and slow planning with language models”, *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- [26] Amith Ananthram, **Elias Stengel-Eskin**, Carl Vondrick, Mohit Bansal, and Kathleen McKeown, “See it from my perspective: Diagnosing the western cultural bias of large vision-language models in image understanding”, *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- [27] **Elias Stengel-Eskin**, Peter Hase, and Mohit Bansal, “Teaching models to balance resisting and accepting persuasion”, *Proceedings of the Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)*, 2025.
- [28] David Wan, Justin Chih-Yao Chen, **Elias Stengel-Eskin**, and Mohit Bansal, “MAMM-Refine: A recipe for improving faithfulness in generation with multi-agent collaboration”, *Proceedings of the Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)*, 2025.
- [29] Han Wang, Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “AdaCAD: Adaptively decoding to balance conflicts between contextual and parametric knowledge”, *Proceedings of the Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)*, 2025.
- [30] **Elias Stengel-Eskin**, Peter Hase, and Mohit Bansal, “LACIE: Listener-aware finetuning for confidence calibration in large language models”, *Advances in Neural Information Processing Systems 38*, 2024.
- [31] Jinhao Duan, Renming Zhang, James Diffenderfer, Bhavya Kailkhura, Lichao Sun, **Elias Stengel-Eskin**, Mohit Bansal, Tianlong Chen, and Kaidi Xu, “GTBench: Uncovering the strategic reasoning limitations of LLMs via game-theoretic evaluations”, *Advances in Neural Information Processing Systems 38*, 2024.

- [32] Ansel Blume, Khanh Duy Nguyen, Zhenhailong Wang, Yangyi Chen, Michal Shlapentokh-Rothman, Xiaomeng Jin, Jeonghwan Kim, ..., **Stengel-Eskin-Elias**, *et al.*, “MIRACLE: An online, explainable multimodal interactive concept learning system”, in *Proceedings of the 32nd ACM International Conference on Multimedia*, 2024.
- [33] Peter Hase, Thomas Hofweber, Xiang Zhou, **Elias Stengel-Eskin**, and Mohit Bansal, “Fundamental problems with model editing: How should rational belief revision work in LLMs?”, *Transactions on Machine Learning*, 2024.
- [34] **Elias Stengel-Eskin**, Archiki Prasad, and Mohit Bansal, “ReGAL: Refactoring programs to discover generalizable abstractions”, *The 41st International Conference on Machine Learning (ICML)*, 2024.
- [35] Justin Chih-Yao Chen, Swarnadeep Saha, **Elias Stengel-Eskin**, and Mohit Bansal, “MAGDi: Structured distillation of multi-agent interaction graphs improves reasoning in smaller language models”, *The 41st International Conference on Machine Learning (ICML)*, 2024.
- [36] Haotian Fu, Pratyusha Sharma, **Elias Stengel-Eskin**, George Konidaris, Nicolas Le Roux, Marc-Alexandre Côté, and Xingdi Yuan, “Language-guided skill learning with temporal variational inference”, *The 41st International Conference on Machine Learning (ICML)*, 2024.
- [37] David Wan, Jaemin Cho, **Elias Stengel-Eskin**, and Mohit Bansal, “Contrastive region guidance: Improving grounding in vision-language models without training”, *The 18th European Conference on Computer Vision (ECCV)*, 2024.
- [38] Han Wang*, Archiki Prasad*, **Elias Stengel-Eskin***, and Mohit Bansal, “Soft self-consistency improves language model agents”, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- [39] Maryam Hashemzadeh, **Elias Stengel-Eskin**, Sarath Chandar, and Marc-Alexandre Cote, “Sub-goal distillation: A method to improve small language agents”, *The 3rd Conference on Lifelong Learning Agents (CoLLAs)*, 2024.
- [40] **Elias Stengel-Eskin**, Kyle Rawlins, and Benjamin Van Durme, “Zero and few-shot semantic parsing with ambiguous inputs”, in *The Twelfth International Conference on Learning Representations (ICLR)*, 2024.
- [41] Archiki Prasad, **Elias Stengel-Eskin**, and Mohit Bansal, “Rephrase, augment, reason: Visual grounding of questions for vision-language models”, in *The Twelfth International Conference on Learning Representations (ICLR)*, 2024.
- [42] **Elias Stengel-Eskin** and Benjamin Van Durme, “Did You Mean...? Confidence-based Trade-offs in Semantic Parsing”, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [43] **Elias Stengel-Eskin** and Benjamin Van Durme, “Calibrated Interpretation: Confidence Estimation in Semantic parsing”, *Transactions of the Association for Computational Linguistics (TACL)*, 2023.
- [44] **Elias Stengel-Eskin**, Jimena Guallar-Blasco, Yi Zhou, and Benjamin Van Durme, “Why Did the Chicken Cross the Road? Rephrasing and Analyzing Ambiguous Questions in VQA”, *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023.
- [45] Zhuowan Li, Xingrui Wang, **Elias Stengel-Eskin**, Adam Kortylewski, Wufei Ma, Benjamin Van Durme, and Alan Yuille, “Super-CLEVR: A Virtual Benchmark to Diagnose Domain Robustness in Visual Reasoning”, *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [46] **Elias Stengel-Eskin** and Benjamin Van Durme, “The Curious Case of Control”, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022.

- [47] **Elias Stengel-Eskin**, Emmanouil Antonios Platanios, Adam Pauls, Sam Thomson, Hao Fang, Benjamin Van Durme, Jason Eisner, and Yu Su, “When More Data Hurts: A Troubling Quirk in Developing Broad-Coverage Natural Language Understanding Systems”, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022.
- [48] Chenyu Zhang, Benjamin Van Durme, Zhuowan Li, and **Elias Stengel-Eskin**, “Visual Commonsense in Pretrained Unimodal and Multimodal Models”, in *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, Seattle, Washington: Association for Computational Linguistics, 2022.
- [49] **Elias Stengel-Eskin**, Andrew Hundt, Zhuohong He, Aditya Murali, Nakul Gopalan, Matthew Gombolay, and Gregory D. Hager, “Guiding Multi-Step Rearrangement Tasks with Natural Language Instructions”, in *5th Annual Conference on Robot Learning (CoRL)*, 2021.
- [50] Zhuowan Li, **Elias Stengel-Eskin**, Yixiao Zhang, Cihang Xie, Quan Tran, Benjamin Van Durme, and Alan Yuille, “Calibrating Concepts and Operations: Towards Symbolic Reasoning on Real Images”, in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021.
- [51] **Elias Stengel-Eskin**, Jimena Guallar-Blasco, and Benjamin Van Durme, “Human-Model Divergence in the Handling of Vagueness”, in *Proceedings of the 1st Workshop on Understanding Implicit and Underspecified Language (UnImplicit)*, Online: Association for Computational Linguistics, 2021.
- [52] **Elias Stengel-Eskin**, Kenton Murray, Sheng Zhang, Aaron Steven White, and Benjamin Van Durme, “Joint Universal Syntactic and Semantic Parsing”, *Transactions of the Association for Computational Linguistics (TACL)*, 2021.
- [53] **Elias Stengel-Eskin**, Jimena Guallar-Blasco, and Benjamin Van Durme, “Exploring Human-Model Divergence Through Vagueness”, *Proceedings of the Society for Computation in Linguistics (SCiL)*, 2021, ***Abstract**.
- [54] Ryan Culkin, J. Edward Hu, **Elias Stengel-Eskin**, Guanghui Qin, and Benjamin Van Durme, “Iterative Paraphrastic Augmentation with Discriminative Span Alignment”, *Transactions of the Association for Computational Linguistics (TACL)*, volume 9, 2021.
- [55] **Elias Stengel-Eskin**, Aaron Steven White, Sheng Zhang, and Benjamin Van Durme, “Universal Decompositional Semantic Parsing”, in *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2020.
- [56] Aaron Steven White, **Elias Stengel-Eskin**, Siddharth Vashishtha, Venkata Subrahmanyam Govindarajan, Dee Ann Reisinger, Tim Vieira, Keisuke Sakaguchi, Sheng Zhang, Francis Ferraro, Rachel Rudinger, *et al.*, “The Universal Decompositional Semantics Dataset and Decomp Toolkit”, in *Proceedings of The 12th Language Resources and Evaluation Conference*, 2020.
- [57] **Elias Stengel-Eskin**, Tzu-Ray Su, Matt Post, and Benjamin Van Durme, “A Discriminative Neural Model for Cross-Lingual Word Alignment”, in *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, 2019.
- [58] Michael McAuliffe, **Elias Stengel-Eskin**, Michaela Socolof, and Morgan Sonderegger, “Polyglot and Speech Corpus Tools: A System for Representing, Integrating, and Querying Speech Corpora.” In *INTERSPEECH*, 2017.

PUBLICATIONS (OTHER)

1. Justin Chih-Yao Chen, Archiki Prasad, Zaid Khan, Joykirat Singh, Runchu Tian, **Elias Stengel-Eskin**, Mohit Bansal, “Cog-DRIFT: Exploration on Adaptively Reformulated Instances Enables Learning from Hard Reasoning Problems” (2026), *Under Review*

2. Marvin Limpijankit, Milad Alshomary, Yassin Oulad Daoud, Amith Ananthram, Tim Trombley, **Elias Stengel-Eskin**, Mohit Bansal, Noam M. Elcott, Kathleen McKeown, “Does AI See like Art Historians? Interpreting How Vision Language Models Recognize Artistic Style” (2026), *Under Review*
3. Joykirat Singh, Justin Chih-Yao Chen, Archiki Prasad, **Elias Stengel-Eskin**, Akshay Nambi, Mohit Bansal, “Think Right: Learning to Mitigate Under-Over Thinking via Adaptive, Attentive Compression” (2025), *Under Review*
4. Vaidehi Patil, **Elias Stengel-Eskin**, Mohit Bansal, “The Sum Leaks More Than Its Parts: Compositional Privacy Risks and Mitigations in Multi-Agent Collaboration” (2025), *Under Review*
5. Hanqi Xiao, Vaidehi Patil, Hyunji Lee, **Elias Stengel-Eskin**, Mohit Bansal, “Generalized Correctness Models: Learning Calibrated and Model-Agnostic Correctness Predictors from Historical Patterns” (2025), *Under Review*
6. Nithin Sivakumaran, Shoubin Yu, Hyunji Lee, Yue Zhang, Ali Payani, Mohit Bansal, **Elias Stengel-Eskin**, “Balancing Faithfulness and Performance in Reasoning via Multi-Listener Soft Execution” (2026), *Under Review*
7. David Wan, Han Wang, Ziyang Wang, **Elias Stengel-Eskin**, Hyunji Lee, Mohit Bansal, “Multimodal Fact-Level Attribution for Verifiable Reasoning” (2026), *Under Review*
8. Duy Nguyen, Hanqi Xiao, Archiki Prasad, **Elias Stengel-Eskin**, Hyunji Lee, Mohit Bansal, “Conflict-Resolving and Sharpness-Aware Minimization for Generalized Knowledge Editing with Multiple Updates” (2026), *Under Review*
9. Haotian Zhai, **Elias Stengel-Eskin**, Pratik Patil, Liu Leqi, “Evaluating Stochasticity in Deep Research Agents” (2026), *Under Review*
10. David Wan, Han Wang, **Elias Stengel-Eskin**, Jaemin Cho, Mohit Bansal, “CLaMR: Contextualized Late-Interaction for Multimodal Content Retrieval” (2025), *Under Review*
11. Emmanouil Zaranis, António Farinhas, Saul Santos, Beatriz Canaverde, Miguel Moura Ramos, Aditya K Surikuchi, André Viveiros, Baohao Liao, Elena Bueno-Benito,..., Nithin Sivakumaran, **Elias Stengel-Eskin**, et al., “Movie Facts and Fibs (MF²): A Benchmark for Long Movie Understanding” (2025), *Under Review*
12. Ron Eliav, Arie Cattan, Eran Hirsch, Shahaf Bassan, **Elias Stengel-Eskin**, Mohit Bansal, Ido Dagan, “CLATTER: Comprehensive Entailment Reasoning for Hallucination Detection” (2025), *Under Review*
13. Justin Chih-Yao Chen, Sukwon Yun, **Elias Stengel-Eskin**, Tianlong Chen, Mohit Bansal, “Symbolic mixture-of-experts: Adaptive skill-based routing for heterogeneous reasoning” (2025), *Under Review*
14. Zaid Khan, **Elias Stengel-Eskin**, Archiki Prasad, Jaemin Cho, Mohit Bansal, “Executable Functional Abstractions: Inferring Generative Programs for Advanced Math Problems” (2025), *Under Review*
15. Vaidehi Patil, **Elias Stengel-Eskin**, Mohit Bansal, “UPCORE: Utility-Preserving Coreset Selection for Balanced Unlearning” (2025), *Under Review*
16. Yue Huang, Chujie Gao, Siyuan Wu, ..., **Elias Stengel-Eskin**, ..., Bo Li, Xiangliang Zhang, “On the Trustworthiness of Generative Foundation Models: Guideline, Assessment, and Perspective” (2025), *Preprint*
17. Thomas Hofweber, Peter Hase, **Elias Stengel-Eskin**, Mohit Bansal, “Are language models rational? The case of coherence norms and belief revision” (2024), *Preprint*
18. **Elias Stengel-Eskin**, “Modeling Meaning for Description and Interaction”, (2023) **Ph.D. Thesis**
19. Jimena Guallar-Blasco, **Elias Stengel-Eskin**, Benjamin Van Durme, “Analyzing Question Ambiguity in Why-Questions” (2023), *Mid-Atlantic Student Colloquium on Speech, Language and Learning (MASC-SLL 2023)* *Abstract
20. Shalaka Vaidya, **Elias Stengel-Eskin**, Joao Sedoc, “Automatic Evaluation of Chit-chat via Semantic Parsing” (2022), *Mid-Atlantic Student Colloquium on Speech, Language and Learning (MASC-SLL 2022)* *Abstract
21. Yunmo Chen, Seth Ebner, Tongfei Chen, Patrick Xia, **Elias Stengel-Eskin**, Tzu-Ray Su, J. Edward Hu, Nils Holzenberger, Ryan Culkin, Craig Harman, Max Thomas, Thomas Lippincott, Aaron Steven White, Kyle Rawlins, Benjamin Van Durme, “NIST TAC SM-KBP 2019 System Description: JHU/UR Framework”, (2019)
22. **Elias Stengel-Eskin**, “Variational Bayesian Inference for Unsupervised Lexicon Discovery” (2017), **Undergraduate Honours Thesis**

FUNDING

Received

- **IARPA BENGAL** 2025-2027
\$290k. Role: Postdoc
- **NSF NAIRR Pilot Grant** 2024-2025
Supervising Calibration for Safe and Trustworthy Artificial Intelligence Systems
2,500 GPU Hours. Role: **PI**
- **Cisco Gift** 2025
Improving LLM Reliability and Trustworthiness via Overconfidence Reduction and Conversation
\$75K. Role: Co-PI
- **CapitalOne Gift** 2025
Multi-Agent Systems and Mixtures-of-Experts for Robust Generation
Role: Co-PI
- **Center for AI Safety Compute Grant** 2024
Listener-Aware Confidence Calibration for Improving LLM Reliability and Trustworthiness
Unlimited A100 GPU access. Role: Co-PI
- **NSF Graduate Research Fellowship** 2018-2023
Role: Student

NEWS COVERAGE

1. *The importance of ambiguity*, Johns Hopkins University CS News, 04/15/2024
2. *Speaking easy: A simple method to clarify voice commands*, Johns Hopkins University CS News, 01/08/2024

INVITED TALKS/PANELS

- **Multi-Model Training for Multi-Agent Communication Skills** 2026
Johns Hopkins University, DSAI Seminar
- **Multi-Model Training for Multi-Agent Communication Skills** 2026
University of Pennsylvania, Clunch
- **Multi-Model Training for Multi-Agent Communication Skills** March 13, 2026
Workshop on Large Language Models, Rice University
- **Panel: Reimagining Work in the Age of AI Agents** March 4, 2026
Texas Symposium on Machine Learning, Responsible AI & Robotics
- **Towards Granular, Robust, Calibrated Judgments via Skill-Based and Persuasion-Balanced Reasoning** December 2, 2025
NeurIPS Expo Workshop on Exploring Trust and Reliability in LLM Evaluation
- **Panel: Agentic AI in Wireless Networking and Communications** November 14, 2025
UT Austin Wireless Networking and Communications Group
- **Panel: The true impact of AI on the labor market** September 30, 2025
Hire EZ
- **Communication, Calibration, and Grounding for Collaborative AI Agents** January 8, 2025
Duke University
- **Confidence-based Rephrasing, Refinement, and Selection** March 22nd, 2024
Keynote – UncertainNLP Workshop, EACL 2024
- **Ambiguity in NLP** October 26th, 2023

- **Uncertainty and Ambiguity in Semantic Structures** October 6th, 2023
Emory University
- **Language, Structure, and Beyond** May 9th, 2023
UNC Chapel Hill – MURGe Group
- **Language, Structure, and Beyond** April 4th, 2023
Scaled Cognition
- **Language, Structure, and Beyond** March 13th, 2023
Brandeis University – CS Department
- **Language, Structure, and Beyond** March 2nd, 2023
University of Cambridge – CS Department
- **Language, Structure, and Beyond** February 17th, 2023
George Mason University – George Mason NLP Group
- **Language, Structure, and Beyond** February 1st, 2023
Georgetown University – NERT Lab
- **Joint Universal Syntactic and Semantic Parsing** April 1st, 2022
Cornell University – Workshop on Meaning in Language

HONORS AND AWARDS

Honors and Awards

- First Class Honours in Cognitive Science 2018
- Dean's Honor List (top 10% of faculty) 2014–2015, 2016–2018
- Arts Undergraduate Research Internship Award 2016
- NSF Graduate Research Fellowship 2019

SERVICE

Reviewing

- **Area Chair:** ARR (2024: 2x, 2025: 4x, 2026: 2x), COLM (2026)
- **Reviewer:** ACL (2021, 2022), NAACL 2022, EACL 2023, CoRL 2023, EMNLP 2023, NeurIPS (2024, 2025), WiNLP (2024, 2025), ICLR (2025, 2026), ICML (2025, 2026), ICCV 2025, TACL (2025, 2026), CVPR 2026, ECCV 2026
- **Program Committee:** UnImplicit 2022, SouthNLP 2024
- **Senior Area Chair:** EMNLP 2026

Organizing

- **Workshop Organizer:** UnImplicit 2024
- **Conference Organizer:** EMNLP 2024 (Publicity Chair), ACL 2026 (Student Research Workshop Chair)
- **Session Chair:** NAACL 2025

PhD/Proposal/RPE Committees at UT Austin

- **Thesis Committees:** Liyan Tang
- **PhD Proposals:** Anibal Solon Heinsfeld, Vanya Cohen, Liyan Tang
- **RPEs:** Fangcong Yin, Meng Chen

External Committees

- **PhD Proposals:** Amith Ananthram (Columbia)

Department Service

- **UT Austin Graduate Admissions Committee**, 2025-2026
- **UT Austin Linguistics Hiring Committee**, 2025-2026
- **CLSP Visit Weekend Committee**, 2018-2022
- **CLSP Admissions Committee**, 2022-2023
- **Application Mentoring Group**, 2022-2023

TEACHING

- **Instructor** at University of Texas at Austin Fall 2025
Communicative and Grounded AI Agents – CS395T
- **Guest Lecturer** at UNC Chapel Hill Fall 2024
Advanced Topics in Natural Language Processing (Aligning LLMs and Humans – COMP 790-158)
- **Guest Lecturer** at UNC Chapel Hill Fall 2024
Deep Learning (COMP 664)
- **Guest Lecturer** at UNC Chapel Hill Fall 2023, Fall 2024
Computational Linguistics
- **Teaching Assistant** at Johns Hopkins University Fall 2019
Artificial Intelligence (EN.601.464/664)

MENTORING

- Victor Wong Summer 2025-Present
Undergraduate Honors Thesis (UT Austin)
- Vaidehi Patil Fall 2024-Present
Ph.D. Student (UNC)
- Zaid Khan Fall 2024-Present
Ph.D. Student (UNC)
- Duy Nguyen Fall 2024-Present
Ph.D. Student (UNC)
- Amith Ananthram Winter 2024-Present
Ph.D. Student (Columbia)
- Ziyang Wang Fall 2023-Fall 2024
Ph.D. Student (UNC)
- Shoubin Yu Fall 2023-Present
Ph.D. Student (UNC)
- David Wan Fall 2023-Present
Ph.D. Student (UNC)
- Justin Chih-Yao Chen Fall 2023-Present
Ph.D. Student (UNC)
- Han Wang Fall 2023-Present
Ph.D. Student (UNC)

• Archiki Prasad <i>Ph.D. Student (UNC)</i>	Summer 2023-Present
• Victor Wang <i>Undergrad (UT Austin)</i> Outstanding Undergraduate Researcher Award, CRA (Honorable Mention)	Summer 2025-Present
• Tianyi Niu <i>Undergrad (UNC)</i>	Summer 2025-Present
• Hanqi Xiao <i>Undergrad (UNC)</i> Outstanding Undergraduate Researcher Award, CRA (Honorable Mention)	Summer 2024-Present
• Atin Pothviraj <i>Undergrad (UNC)</i>	Summer 2024-Present
• Nithin Sivakumaran <i>Undergrad (UNC)</i> Outstanding Undergraduate Researcher Award, CRA (Honorable Mention)	Summer 2024-Present
• Bryan Sukidi <i>Undergrad (UNC)</i>	Fall 2024-Fall 2025
• Evan Paces-Wiles <i>Undergrad (UNC)</i>	Fall 2024-Present
• Jio Choi <i>Undergrad (UNC)</i>	Fall 2024-Summer 2025
• Jimena Guallar-Blasco <i>Undergrad/MS 2024 (JHU, Current: Charles River Analytics)</i> Outstanding Undergraduate Researcher Award, CRA (Honorable Mention)	Summer 2020-Summer 2024
• Maryam Hashemzadeh <i>Ph.D. Student (MILA)</i>	Summer 2023
• Zhuohong (Zooey) He <i>MSE, Spring 2021 (JHU, Current: Intuitive)</i>	Winter 2021-Spring 2021
• Chenyu (Heidi) Zhang <i>BS 2022, (Next: MS (Stanford), Current: Google Gemini)</i>	Fall 2021-Fall 2022
• Shalaka Vaidya <i>MS 2023 (NYU)</i>	Fall 2022
• Yi Zhou <i>MS 2022 (JHU)</i>	Winter 2022-Fall 2022
• Vicky Zeng <i>Ph.D. Student (JHU)</i>	Winter 2021-Winter 2022

SKILLS

- **Programming (expert):** Python
- **Programming (proficient):** Bash, Java, Javascript
- **Libraries/Frameworks:** PyTorch, AllenNLP, Transformers, NLTK, numpy, MechanicalTurk, networkx, MXNet, React

LANGUAGES

- **Native:** English, German
- **Fluent:** French
- **Other:** Latin (reading/translation), Spanish (intermediate)